

GPU Computing

Patterns for massively parallel programming (part 1)

The Global Memory

Edwin Carlinet, Joseph Chazalon

`firstname.lastname@epita.fr`

Spring 2024

EPITA Research Laboratory (LRE)



```
!include agenda.inc.md
```

Programming patterns & Memory Optimizations

The Programming Patterns

- Map
- Map + Local reduction
- Reduction
- Scan

The IP algorithms

- LUT Application
- Local features extraction
- Histogram
- Integral images


Map Pattern

Map Pattern Overview

Map replicates a function over every element of an index set

The computation of each pixel is independent w.r.t. the others.

$\text{out}(x,y) = f(\text{in}(x,y))$



Nothing complicated but take care of memory access pattern.

```
void plus_one(int* a, int size, int k) {  
    int i = blockDim.x * blockIdx.x +  
           threadIdx.x + k;  
    if (i < size)  
        a[i] = a[i] + 1;  
}
```



```
void plus_one(int* a, int size, int k) {  
    int i = blockDim.x * blockIdx.x +  
           threadIdx.x * k;  
    if (i < size)  
        a[i] = a[i] + 1;  
}
```

What do you think about k's impact of the performance?

```
void plus_one(int* a, int size, int k) {
    int i = blockDim.x * blockIdx.x +
            threadIdx.x + k;
    if (i < size)
        a[i] = a[i] + 1;
}
```

```
void plus_one(int* a, int size, int k) {
    int i = blockDim.x * blockIdx.x +
            threadIdx.x * k;
    if (i < size)
        a[i] = a[i] + 1;
}
```


What do you think about k's impact of the performance?




- Linear sequential access with offset (left) → 👍
- Strided access → 🙄

Strided access pattern

- $k = 1$: 6.4 ms (with 600K values)



- $k = 51$: 28.7 ms



The Global Memory 🌶️🌶️

Memory Bandwidth

What you think about memory



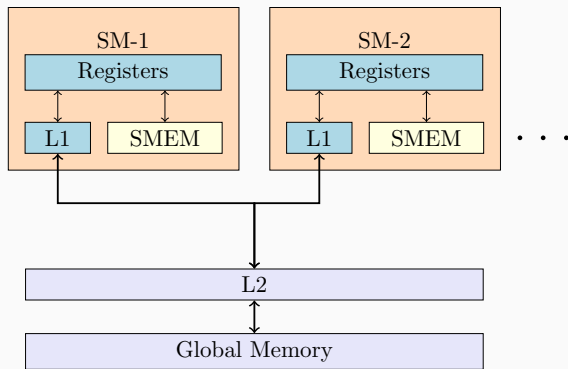
Memory Bandwidth

What you think about memory



Reality

Memory Access Hierarchy



GTX 1080 (Pascal)		Size	Bandwidth	Bus interface	Latency
L1 Cache (per SM)	Low latency	16 or 48K	1,600 GB/s	128 bits	10-20 cycles
L2 Cache		1-2M			
Global	High latency	8GB	320 GB/s	256~384 bits	400-800 cycles

Cached Loads from L1 low-latency memory (1/2)

Cached vs Un-cached

Two types of global memory loads: **Cached** (by default) or **Uncached** (L1 disabled)

Cached Loads from L1 low-latency memory (1/2)

Cached vs Un-cached

Two types of global memory loads: **Cached** (by default) or **Uncached** (L1 disabled)

Aligned vs Misaligned

A load is **aligned** if the first address of a memory access is multiple of 32 bytes

- Memory addresses must be type-aligned (ie `sizeof(T)`)
- Otherwise: poor perf (unaligned load)
- `cudaMalloc` = alignment on 256 bits (at least)

Cached Loads from L1 low-latency memory (1/2)

Cached vs Un-cached

Two types of global memory loads: **Cached** (by default) or **Uncached** (L1 disabled)

Aligned vs Misaligned

A load is **aligned** if the first address of a memory access is multiple of 32 bytes

- Memory addresses must be type-aligned (ie `sizeof(T)`)
- Otherwise: poor perf (unaligned load)
- `cudaMalloc` = alignment on 256 bits (at least)

Coalesced versus uncoalesced


A load is **coalesced** if a warp accesses a contiguous chunk of data

- Minimize memory accesses caused by warp threads
- Remind: all threads of a warp executes the same instruction
→ if a load, may be 32 different addresses

Cached Loads from L1 low-latency memory (2/2)

We need a load strategy:


- 32 threads of warp access a 32-bit word = 128 bytes
- 128 bytes = L1 bus width (single load - bus utilization = 100%)
- Access permutation has no (or very low) overhead



Misaligned cached loads from L1

- If data are not 128-bits aligned, two loads are required


Adresses 96-224 required... but 0-256 loaded



Misaligned cached loads from L1


- If data are not 128-bits aligned, two loads are required

Addresses 96-224 required... but 0-256 loaded



-
- If data is accessed strided


e.g. $u[2*k], \dots$



Misaligned cached loads from L1


- If data are not 128-bits aligned, two loads are required

Addresses 96-224 required... but 0-256 loaded




- If data is accessed strided

e.g. $u[2*k], \dots$




...cry !




Loads from global (uncached) memory 🌶️🌶️

Same idea but memory is split in segments of 32 bytes


Cached



Uncached



Aligned accesses (sequential/non-sequential)



Compute capability:

2.0 and later

Memory transactions:


Uncached

Cached

1x 32B at 128
1x 32B at 160
1x 32B at 192
1x 32B at 224

1x 128B at 128

Mis-aligned accesses (sequential/non-sequential)



Compute capability:

2.0 and later

Memory transactions:

Uncached

Cached

1x 32B at 128
1x 32B at 160
1x 32B at 192
1x 32B at 224
1x 32B at 256

1x 128B at 128
1x 128B at 256

Why would you need uncached memory ? 🌶️🌶️🌶️

Caching

- Better performance if non-coalesced access and data-reuse

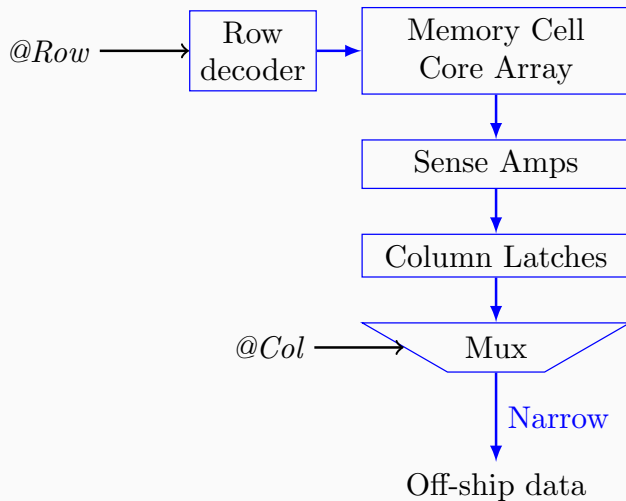
Non-caching

- Avoid wasting cache for one-time used data (stream usage) (→ more space for register spilling)

Memory architecture 🌶️🌶️🌶️


Why cacheline ? Why Aligment ?

- DRAM is organized in 2D Core arrays
- Each DRAM core array has about 16M bits



Example

- A 4 x 4 memory cell
- With 4 bits pin interface width



Reading from a cell in the core array is a very slow process ($1/N$ th of the interface speed):

- DDR: Core speed = $1/2$ interface speed
- DDR2/GDDR3: Core speed = $1/4$ interface speed
- DDR3/GDDR4: Core speed = $1/8$ interface speed

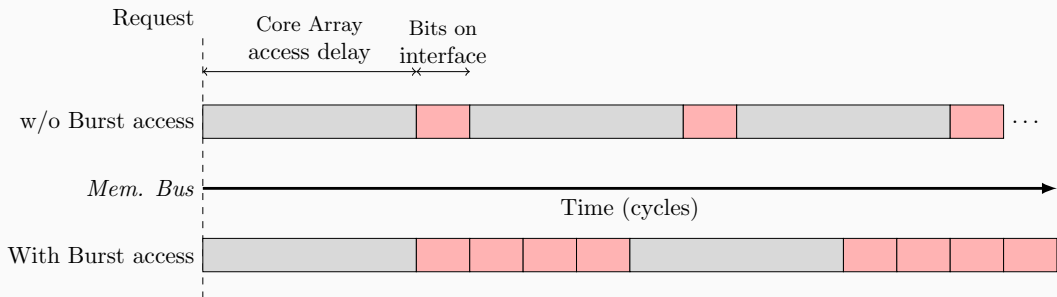
DRAM Burst

Reading from a cell in the core array is a very slow process ($1/N$ th of the interface speed):

- DDR: Core speed = $1/2$ interface speed
- DDR2/GDDR3: Core speed = $1/4$ interface speed
- DDR3/GDDR4: Core speed = $1/8$ interface speed


Solution: Bursting

Load $N \times$ interface width of **the same row**



Better, but not enough to saturate the memory bus 🙅 (will see later a solution).

- Use **coalesced** (*contiguous and aligned*) access to memory:



- If all threads of a warp execute a load instruction into the same burst section → only one DRAM request 👍
- Otherwise:
 - Multiple DRAM requests are made
 - Some bytes transferred are not used


Q: how to make coalesced/aligned loads with 2D-arrays ? 🌶️

M 0,0	M 1,0	M 2,0	M 3,0	M 4,0
M 0,1	M 1,1	M 2,1	M 3,1	M 4,1
M 0,2	M 1,2	M 2,2	M 3,2	M 4,2
M 0,3	M 1,3	M 2,3	M 3,3	M 4,3

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19
Burst section				Burst section				Burst section				Burst section				Burst section			

Q: how to make coalesced/aligned loads with 2D-arrays ? 🌶️

M 0,0	M 1,0	M 2,0	M 3,0	M 4,0
M 0,1	M 1,1	M 2,1	M 3,1	M 4,1
M 0,2	M 1,2	M 2,2	M 3,2	M 4,2
M 0,3	M 1,3	M 2,3	M 3,3	M 4,3



1. Add padding to align rows on 256-bits boundaries
2. Thread reads data column-wise

`ima(tid.x,tid.y) = tid.y * pitch + tid.x`

