

# MLRF Lecture 02


J. Chazalon, LRE/EPITA, 2025

# Local feature detectors

Lecture 02 part 02

# Introduction

*How are panorama pictures created from multiple pictures?*



# Introduction


*How are panorama pictures created from multiple pictures?*



1. Detect small parts invariant under viewpoint change: “Keypoints”

# Introduction


*How are panorama pictures created from multiple pictures?*



1. Detect small parts invariant under viewpoint change: **keypoints**
2. Find pairs of matching keypoints using a **description** of their neighborhood

# Introduction

*How are panorama pictures created from multiple pictures?*



1. Detect small parts invariant under viewpoint change: **keypoints**
2. Find pairs of matching keypoints using a **description** of their neighborhood
3. Compute the **most likely transformation** to blend images together

# The need for local feature detectors

While **dense computation** of local feature descriptors is possible (grid of points), this is **rarely used in practice** (lots of computations, lots of useless features).

**Detection** = Find **anchors** to describe a **feature of interest**.

- Edge / line
- Area around a corner / a stable point
- Blob (area of variable size)

A good feature of interest is **stable over the perturbations** our signal will face:

- Translation, rotation, zoom, perspective
- Illumination changes
- Noise, compression
- ...

# Some classical detectors

## Edge (gradient detectors)

- Sobel
- Canny

## Corner

- Harris & Stephens *and variants*
- FAST
- Laplacian of Gaussian, Difference of Gaussian, Determinant of Hessian

## Blob

- MSER



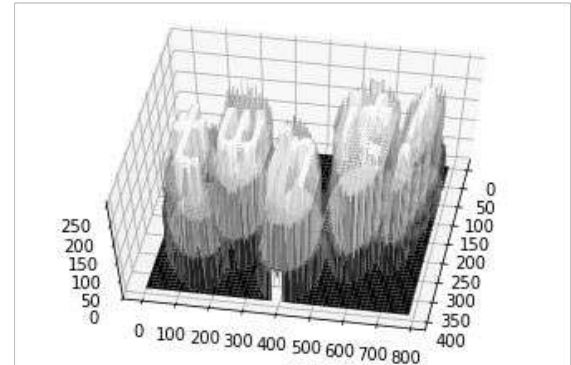



# Edge detectors

# What's an edge?

Image is a function

Edges are rapid changes in this function

The derivative of a function exhibits the edges



# Image derivatives

Recall: 
$$f'(a) = \lim_{h \rightarrow 0} \frac{f(a+h) - f(a-h)}{2h}$$


We don't have an “actual” function, must estimate

Possibility: set  $h = 1$

Apply filter 

-1	0	+1
----	---	----

 to the image  
(x gradient)







# Image derivatives

We get terribly spiky results,  
we need to interpolate / smooth.

⇒ Gaussian filter




We get a Sobel filter



$$\frac{1}{2} \times \left( \begin{bmatrix} -1 & 0 & +1 \end{bmatrix} * \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} \right) = \begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix} \quad \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

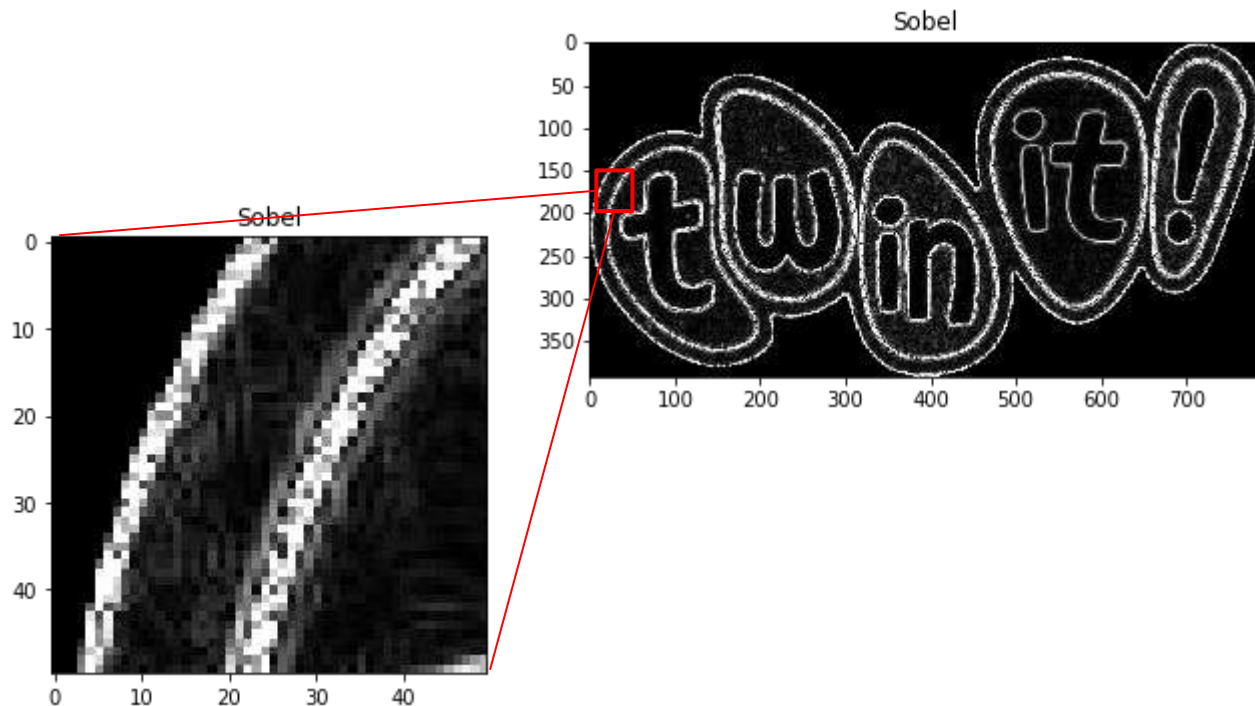
Horizontal Sobel      Vertical Sobel

# Sobel filter



# Gradient magnitude with Sobel

$\text{sqrt}(\text{Sobel\_x}^2 + \text{Sobel\_y}^2)$



# Canny edge detection


Extract real lines!

Algorithm:

Sobel operator


- Smooth image (only want “real” edges, not noise)
- Calculate gradient direction and magnitude
- Non-maximum suppression perpendicular to edge
- Threshold into strong, weak, no edge
- Keep only weak pixels connected to strong ones

# Canny: Non-maximum suppression







# Canny: Non-maximum suppression




# Canny: Non-maximum suppression




# Canny: Non-maximum suppression




# Canny: Non-maximum suppression




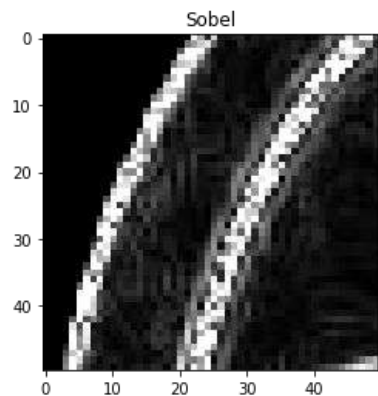
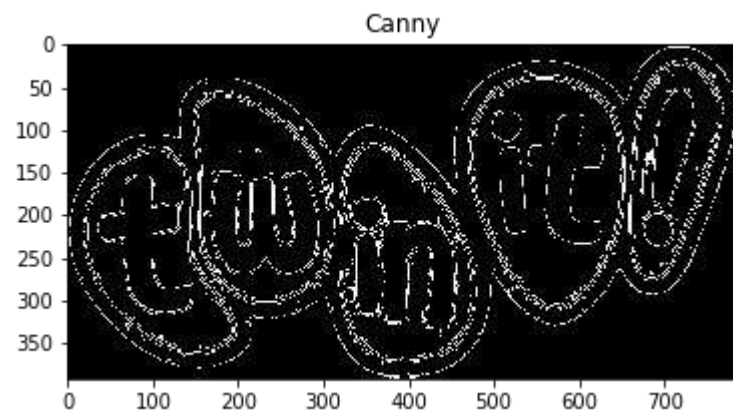

# Canny: Non-maximum suppression



# Canny: Non-maximum suppression



# Canny: Non-maximum suppression



# Canny: finalization

## Threshold edges

- Still some noise
- Only want strong edges
- 2 thresholds, 3 cases
  - $R > T$ : strong edge
  - $R < T$  but  $R > t$ : weak edge
  - $R < t$ : no edge
- Why two thresholds?  
*“Hysteresis thresholding”*

## Connect weak edges to strong edges

- Strong edges are edges!
- Weak edges are edges  
iff they connect to strong
- Look in some neighborhood  
(usually 8 closest)



# Corner detectors

## Introduction, Harris detector

# Good features

Reminder:

Good features are unique!

- Can find the “same” feature easily
- Not mistaken for “different” features

Good features are robust under perturbation

- Can detect them under translation, rotation...
- Intensity shift...
- Noise...

How close are two patches?

- Sum squared difference
- Images I, J
- $\sum_{x,y} (I(x,y) - J(x,y))^2$

# How can we find unique patches?

Say we are stitching a panorama

Want patches in image to match to other image


Need to only match one spot



# How can we find unique patches?

## Sky? Bad!

- Very little variation
- Could match any other sky




# How can we find unique patches?

## Sky? Bad!

- Very little variation
- Could match any other sky

## Edge? OK...

- Variation in one direction
- Could match other patches along same edge



# How can we find unique patches?

## Sky? Bad!


- Very little variation
- Could match any other sky

## Edge? OK...

- Variation in one direction
- Could match other patches along same edge

## Corners? good!

- Only one alignment matches



# How can we find unique patches?

Want a patch that is unique in the image

Can calculate distance between patch  
and every other patch, lot of computation



# How can we find unique patches?

Want a patch that is unique in the image

Can calculate distance between patch  
and every other patch, lot of computation

Instead, we could think about  
auto-correlation:

How well does an image match  
a shifted version of itself?



$$\sum_{\mathbf{d}} \sum_{x,y} (I(x,y) - I(x+\mathbf{d}_x, y+\mathbf{d}_y))^2$$

Measure of self-difference  
(how am I not myself?)





# Self-difference

Sky: low everywhere





# Self-difference

Edge: low along edge




# Self-difference

Corner: mostly high




# Self-difference


Corner: mostly high



Edge: low along edge




Sky: low everywhere



# Self-difference

Naive computation:

$$\sum_d \sum_{x,y} (I(x,y) - I(x+d_x, y+d_y))^2$$




$$(I(x,y) - I(x+d_x, y+d_y))^2$$


# Harris corner detector

In practice we pool the previous indicator function over a small region  $(u,v)$  and we use a window  $w(u,v)$  to weight the contribution of each displacement to the global sum.


$$S(x, y) = \sum_u \sum_v w(u, v) \left( I(x + u + d_x, y + v + d_y) - I(x + u, y + v) \right)^2$$



$$(I(x, y) - I(x + \mathbf{d}_x, y + \mathbf{d}_y))^2$$



$$\sum_u \sum_v$$



$$w(u, v)$$




# Harris corner detector

$$\sum_d \sum_{x,y} (I(x,y) - I(x+d_x, y+d_y))^2$$

Lots of summing => Need an approximation

Look at nearby gradients  $I_x$  and  $I_y$

- If gradients are **mostly zero**, not a lot going on  
⇒ Low self-difference
- If gradients are **mostly in one direction**, edge  
⇒ Still low self-difference
- If gradients are **in twoish directions**, corner!  
⇒ High self-difference, good patch!



# Harris corner detector

Trick to precompute the derivatives

$$I(x + d_x, y + d_y)$$

can be approximated by a Taylor expansion

$$I(x + d_x, y + d_y) \approx I(x, y) + d_x \frac{\partial I(x, y)}{\partial x} + d_y \frac{\partial I(x, y)}{\partial y} + \dots$$



# Harris corner detector

This allows us to "simplify" the original equation,

$$S(x, y) \approx \sum_u \sum_v w(u, v) \left( d_x \frac{\partial I(x + u, y + v)}{\partial x} + d_y \frac{\partial I(x + u, y + v)}{\partial y} \right)^2$$

and more important making it **faster to compute**,  
thanks to simpler derivatives which can be **computed for the whole image**.

# Harris corner detector

If we develop the equation and write it as usual matrix form, we get:

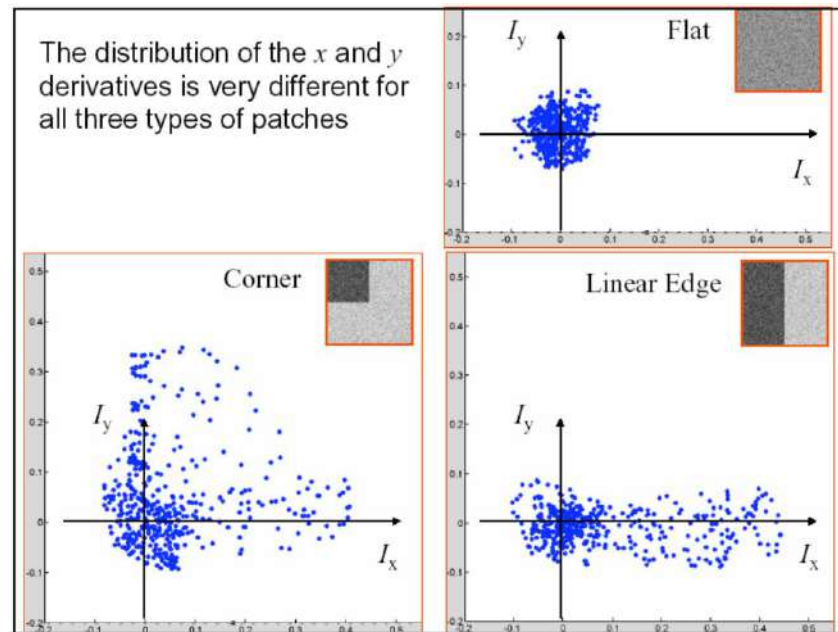

$$S(x, y) \approx \begin{pmatrix} d_x & d_y \end{pmatrix} A(x, y) \begin{pmatrix} d_x \\ d_y \end{pmatrix}$$

where  $A(x, y)$  is the structure tensor:


$$A = \sum_u \sum_v w(u, v) \begin{bmatrix} \frac{\partial^2 I(x+u, y+v)}{\partial x^2} & \frac{\partial I(x+u, y+v)}{\partial x} \frac{\partial I(x+u, y+v)}{\partial y} \\ \frac{\partial I(x+u, y+v)}{\partial x} \frac{\partial I(x+u, y+v)}{\partial y} & \frac{\partial^2 I(x+u, y+v)}{\partial y^2} \end{bmatrix}$$
$$= \begin{bmatrix} \langle I_x^2 \rangle & \langle I_x I_y \rangle \\ \langle I_x I_y \rangle & \langle I_y^2 \rangle \end{bmatrix}$$

This trick is useful because  $I_x$  and  $I_y$  can be precomputed very simply.

# Harris corner detector




# Harris corner detector



The need for eigenvalues:

If the edge is rotated,  
so are the values of  $I_x$  and  $I_y$ .

Eigenvalues give us the ellipsis axis len.




# Harris corner detector

A corner is characterized by a large variation of  $S$  in all directions of the vector  $(x, y)$ .

Analyse the eigenvalues of  $A$  to check whether we have two large variations.

- If  $\lambda_1 \approx 0$  and  $\lambda_2 \approx 0$  then this pixel  $(x, y)$  has no features of interest.
- If  $\lambda_1 \approx 0$  and  $\lambda_2$  has some large positive value, then an edge is found.
- If  $\lambda_1$  and  $\lambda_2$  have large positive values, then a corner is found.



# Harris corner detector

To avoid the computation of the eigenvalues, which used to be expensive, Harris and Stephens instead suggest the following function  $M_c$ , where  $\kappa$  is a tunable sensitivity parameter:

$$M_c = \lambda_1 \lambda_2 - \kappa (\lambda_1 + \lambda_2)^2 = \underbrace{\det(A) - \kappa \text{trace}^2(A)}_{\text{approximation}}$$

We will use Noble's trick to remove  $\kappa$ :

$$M'_c = 2 \frac{\det(A)}{\text{trace}(A) + \epsilon}$$

$\epsilon$  being a small positive constant.

# Harris corner detector

$A$  being a 2x2 matrix, we have the following relations:

- $\det(A) = A_{1,1}A_{2,2} - A_{2,1}A_{1,2}$
- $\text{trace}(A) = A_{1,1} + A_{2,2}$

Using previous definitions, we obtain:

- $\det(A) = \langle I^2_x \rangle \langle I^2_y \rangle - \langle I_x I_y \rangle^2$
- $\text{trace}(A) = \langle I^2_x \rangle + \langle I^2_y \rangle$

# Harris corner detector

In summary, given an image, we can compute the Harris corner response image by simply computing:


- $I_x$  :  $I$ 's smoothed (interpolated) partial derivative with respect to  $x$  ;
- $I_y$  :  $I$ 's smoothed (interpolated) partial derivative with respect to  $y$  ;
- $\langle I^2_x \rangle$  : the windowed sum of  $I^2_x$  ;
- $\langle I^2_y \rangle$  : the windowed sum of  $I^2_y$  ;
- $\langle I_x I_y \rangle$  : the windowed sum of  $I_x I_y$  ;
- $\det(A)$  ;
- $\text{trace}(A)$  ;
- $M''_c = \det(A) / (\text{trace}(A) + \epsilon)$ .

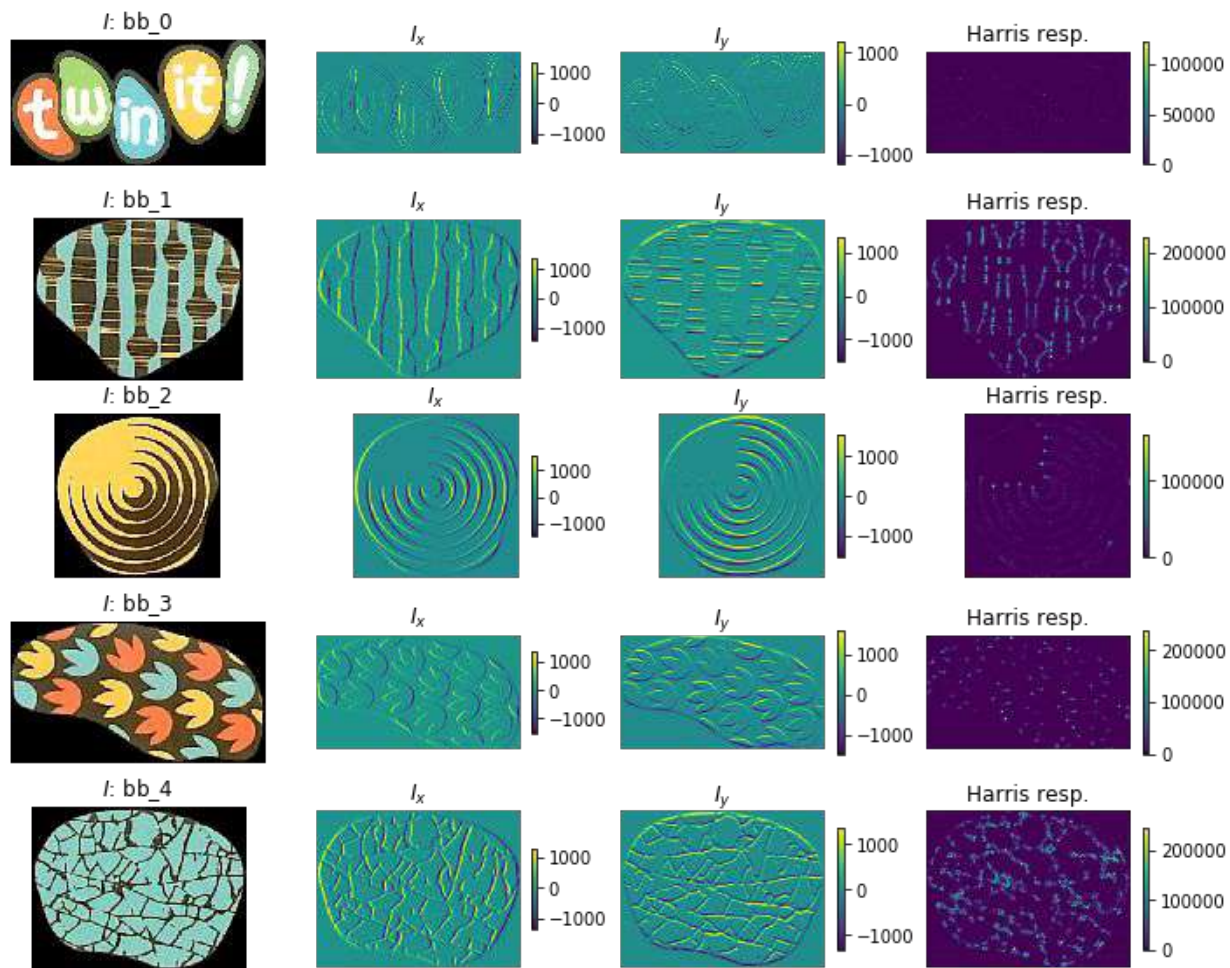
Then, we just perform **non-maximal suppression** to keep local maximas.





















# Harris & Stephens Conclusion

# Good features to track *aka* Shi-Tomasi *aka* Kanade-Tomasi

Remember the Harris-Stephens trick to avoid computing the eigenvalues?

$$M_c = \lambda_1 \lambda_2 - \kappa (\lambda_1 + \lambda_2)^2 = \underbrace{\det(A) - \kappa \text{trace}^2(A)}_{\text{approximation}}$$

Well, nowadays, linear algebra is cheap, so **compute the real eigenvalues**.

Then filter using  $\min(\lambda_1, \lambda_2) > \lambda$ ,  $\lambda$  being a predefined threshold.

You get the Shi-Tomasi variant.

# Build your own edge/corner detector

Hessian matrix with  
block-wise summing

You just need eigenvalues  $\lambda_1$  and  $\lambda_2$  of the structure tensor

$$A = \sum_u \sum_v w(u, v) \begin{bmatrix} \frac{\partial^2 I(x+u, y+v)}{\partial x^2} & \frac{\partial I(x+u, y+v)}{\partial x} \frac{\partial I(x+u, y+v)}{\partial y} \\ \frac{\partial I(x+u, y+v)}{\partial x} \frac{\partial I(x+u, y+v)}{\partial y} & \frac{\partial^2 I(x+u, y+v)}{\partial y^2} \end{bmatrix}$$
$$= \begin{bmatrix} \langle I_x^2 \rangle & \langle I_x I_y \rangle \\ \langle I_x I_y \rangle & \langle I_y^2 \rangle \end{bmatrix}$$

```
dst = cv2.cornerEigenValsAndVecs(src, neighborhood_size, sobel_aperture)
dst = cv2.cornerMinEigenVal(src, neighborhood_size, sobel_aperture)
```



# Harris summary

## Pros

Translation invariant

⇒ Large gradients in both directions  
= stable point

## Cons

**Not** so fast

⇒ Avoid to compute all those derivatives

**Not** scale invariant

⇒ Detect corners at different *scales*

**Not** rotation invariant

⇒ Normalization orientation



# Corner detectors, binary tests

## FAST

# Features from accelerated segment test (FAST)

*Keypoint detector used by ORB (described in next lecture)*


## Segment test:

compare pixel  $P$  intensity  $I_p$   
with surrounding pixels  
(circle of 16 pixels)

If  $n$  contiguous pixels are either

- all darker than  $I_p - t$
- all brighter than  $I_p + t$

then  $P$  is detected as a corner



**Figure 1.** 12 point segment test corner detection in an image patch. The highlighted squares are the pixels used in the corner detection. The pixel at  $p$  is the centre of a candidate corner. The arc is indicated by the dashed line passes through 12 contiguous pixels which are brighter than  $p$  by more than the threshold.

# Tricks

1. **Cascading:** If  $n = 12$  ( $\frac{3}{4}$  of the circle), then many non-corners can be discarded by testing pixels at the 4 compass directions. The full test is only applied to the candidates which passed the first test.
2. **Machine learning:** Learn on a dataset which pixels should be tested first to discard a non-corner as quickly as possible.

*Learn a decision tree, then compile the decisions as nested if-then rules.*

3. How to perform **non-maximal suppression**?

Need to assign a score  $V$  to each corner.

⇒ The sum of the absolute difference between the pixels in the contiguous arc and the centre pixel

$$V = \max \left( \sum_{x \in S_{\text{bright}}} |I_{p \rightarrow x} - I_p| - t, \sum_{x \in S_{\text{dark}}} |I_p - I_{p \rightarrow x}| - t \right)$$

# FAST summary

## Pros

Very fast

*Authors tests:*

- 20 times faster than Harris
- 40 times faster than DoG (*next slide*)

Very robust to transformations (perspective in particular)

## Cons

Very sensitive to blur

# Corner detectors at different scales

## LoG, DoG, DoH

# Laplacian of Gaussian (LoG)

The theoretical, slow way.


If you need to remember only 1 thing:

it is a **band-pass filter** – it **detects objects of a certain size**.

*original image*



*scale-space maxima of  $(\nabla_{norm}^2 L)^2$*



# Laplacian (plain, not Gaussian here) = second derivative

Second derivative of an image? Like Sobel... with 1 more derivation...

Taylor, again:

$$f(x+h) = f(x) + hf'(x) + \frac{1}{2}h^2 f''(x) + \frac{1}{3!}h^3 f'''(x) + O(h^4)$$

add

$$+ \left[ f(x-h) = f(x) - hf'(x) + \frac{1}{2}h^2 f''(x) - \frac{1}{3!}h^3 f'''(x) + O(h^4) \right]$$

---

$$f(x+h) + f(x-h) = 2f(x) + h^2 f''(x) + O(h^4)$$

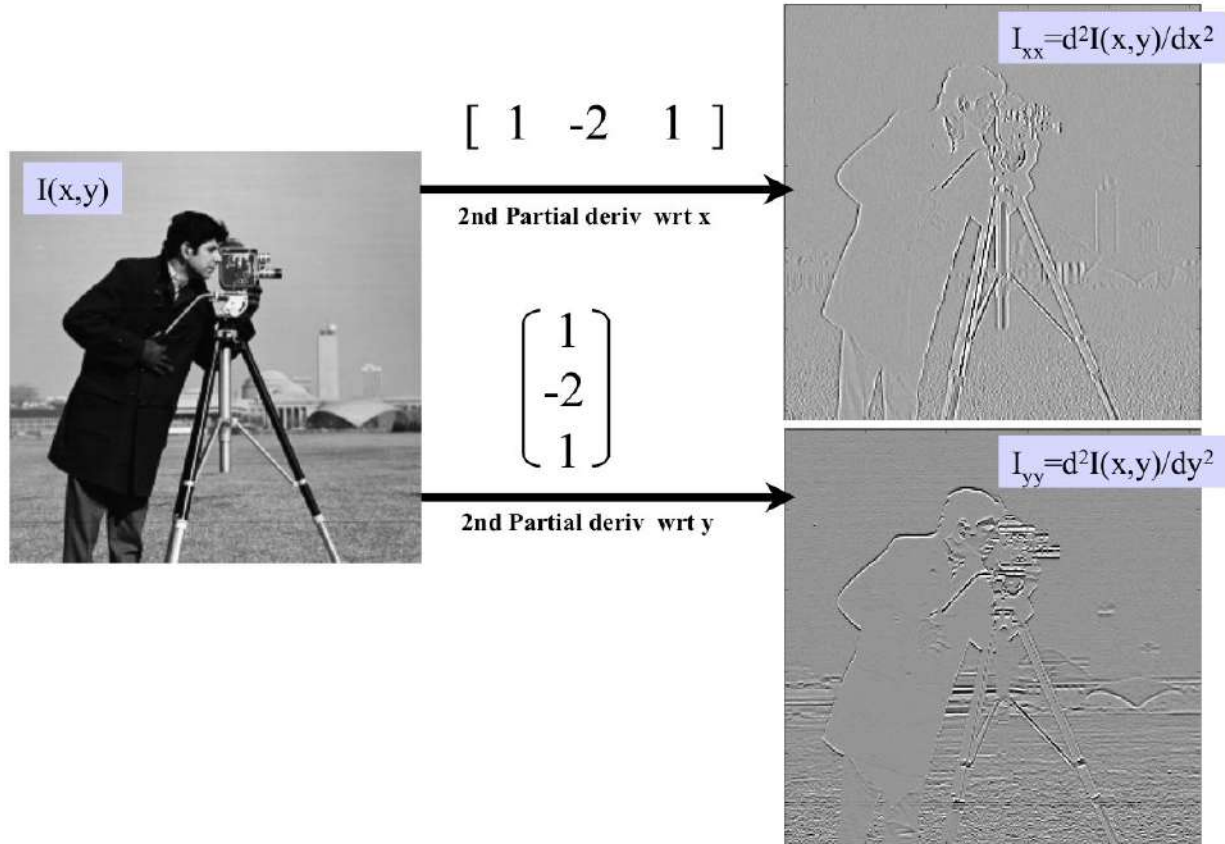
$$\frac{f(x-h) - 2f(x) + f(x+h)}{h^2} = f''(x) + O(h^2)$$

New filter:  $I_{xx} =$ 

1	-2	1
---	----	---

 $\ast /$

# Second partial derivatives of an image






# Laplacian filter $\nabla^2 I(x,y)$

Edge detector, like Sobel but with 2nd derivatives

$$I_{xx} + I_{yy} = \left( \begin{bmatrix} 1 & -2 & 1 \end{bmatrix} + \begin{bmatrix} 1 \\ -2 \\ 1 \end{bmatrix} \right) * I$$
$$= \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} * I$$



$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} * I$$




# Laplacian of Gaussian




Second derivative of a Gaussian: “**Mexican hat**”

$$g''(x) = \left( \frac{x^2}{\sigma^4} - \frac{1}{\sigma^2} \right) e^{-\frac{x^2}{2\sigma^2}}$$

*2D formula = exercise.*



# LoG = detector of circular shapes




# LoG = detector of circular shapes

LoG filter extrema locates “blobs”


- maxima = dark blobs on light background
- minima = light blobs on dark background

**Scale** of blob (size ; radius in pixels) is determined by the **sigma** parameter of the LoG filter.

LoG  $\sigma=2$





LoG  $\sigma=10$



# Detecting corners / blobs

Build a scale space representation: *stack of images (3D) with increasing sigma*



Then find **local extremas** in the scale space volume.

# Difference of Gaussian (DoG)



Fast approximation of LoG. Used by SIFT (next part of the lecture).

LoG can be approximate by a Difference of two Gaussians (DoG) at different scales.

$$\nabla^2 G_\sigma \approx G_{\sigma_1} - G_{\sigma_2}$$

Best approximation when:

$$\sigma_1 = \frac{\sigma}{\sqrt{2}}, \sigma_2 = \sqrt{2}\sigma$$




# DoG filter

It is a band-pass filter.

$$\Gamma_{\sigma, K\sigma}(x, y) = I * \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/(2\sigma^2)} - I * \frac{1}{2\pi K^2\sigma^2} e^{-(x^2+y^2)/(2K^2\sigma^2)}$$


$$\Gamma_{\sigma, K\sigma}(x, y) = I * \left( \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/(2\sigma^2)} - \frac{1}{2\pi K^2\sigma^2} e^{-(x^2+y^2)/(2K^2\sigma^2)} \right)$$



# DoG filter

## Intuition

- Gaussian ( $g$ ) is a low pass filter
- Strongly reduce components with frequency  $f < \sigma$
- $(g * I)$  low frequency components
- $I - (g * I)$  high frequency components
- $g(\sigma_1) * I - g(\sigma_2) * I \Leftarrow$  Components in between these frequencies
- $g(\sigma_1) * I - g(\sigma_2) * I = [g(\sigma_1) - g(\sigma_2)] * I$





# DoG computation in practice

Take a image.



Blur it.




Take the difference.



# Difference-of-Gaussian filter

Many applications.

Indicates the “size” of the “stable” region around a pixel at a given freq. band.





# DoG scale generation trick

## DoG computation: use “octaves”

- “Octave” because frequency doubles/halves between octaves
- If  $\sigma = \sqrt{2}$ , then 3 levels per octave
- Downsample images for next octave (like double sized kernel)
- Compute the DoG between images

Illustration: D. Lowe



# DoG: Corner selection

Throw out weak responses and edges

Estimate gradients

- Similar to Harris, look at nearby responses
- Not whole image, only a few points! Faster!
- Throw out weak responses

Find cornery things

- Same deal, structure matrix, use det and trace information (SIFT variant)

→  
D. G. Lowe, "Distinctive image features from  
scale-invariant keypoints," *International  
journal of computer vision*, vol. 60, no. 2, pp.  
91–110, 2004., see p. 12

$$\frac{\text{Tr}(\mathbf{H})^2}{\text{Det}(\mathbf{H})} < \frac{(r + 1)^2}{r} \quad \mathbf{H} = \begin{bmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{bmatrix}$$

76

# Determinant of Hessian (DoH)

Faster approximation. Used by SURF.

Better resistance to perspective

Computes the scale-normalized determinant of the Hessian (strength of the curvature at a given point)

$$\det H_{norm} L = \sigma^2 (L_{xx} L_{yy} - L_{xy}^2)$$

⇒ Precompute  $L_{xx}$ ,  $L_{yy}$ ,  $L_{xy}$

⇒ Blur them with the right sigma while computing **det  $H$   $L$** : 3 additions

⇒ normalize: different scales – same value range



$\nabla^2 L$



$\det \mathcal{H} L$






Illustration: T. Lindeberg

local extrema of  $\nabla^2 L$




local extrema of  $\det \mathcal{H} L$




# LoG vs DoG vs DoH


Laplacian of Gaussian



Difference of Gaussian



Determinant of Hessian



# LoG, DoG, DoH summary

## Pros

Very robust to transformations

- Perspective
- Blur

Adjustable size (scale)

## Cons

Slow

# Blob detectors

## MSER



# Maximally Stable Extremal Regions (MSER)

*Detects regions which are stable over thresholds.*

## 1. Create min- & max-tree of the image

tree of included components  
when thresholding the image  
at each possible level







image  $f$

$\equiv$




max-tree  $\mathcal{T}_{\max}(f)$



# Maximally Stable Extremal Regions (MSER)

2. **Select most stable regions** between  $t-\Delta$  and  $t+\Delta$

$R_{t^*}$  is maximally stable iff  $q(t) = |R_{t-\Delta} \setminus R_{t+\Delta}| / |R_t|$   
has local minimum at  $t^*$



$|R| = \text{card}(R)$ ;  $\Delta = \text{parameter}$ ;  $R_{t-\Delta} \setminus R_{t+\Delta} = \text{set difference}$





image  $f$

$\equiv$



max-tree  $\mathcal{T}_{\max}(f)$

# MSER summary

## Pros

Very robust to transformations

- Affine transformations
- Intensity changes

Quite fast

## Cons

Not robust to blur

# Local feature detectors

## Conclusion

# Local feature detectors: Conclusion

Harris Stephens: Can be very stable when combined with DoG

Shi-Tomasi: Assumes affine transformation (avoid it with perspective)

DoG: slow but very robust (perspective, blur, illumination)

DoH: faster than DoG, misses small elements, better with perspective.

FAST: very fast, robust to perspective change (like DoG), but blur quickly kills it

MSER: fast, very stable, good choice when no blur

# Classification

Feature detector	<u>Edge</u>	<u>Corner</u>	<u>Blob</u>
<u>Canny</u>	X		
<u>Sobel</u>	X		
<u>Harris &amp; Stephens / Plessey / Shi-Tomasi</u>	X	X	
<u>Shi &amp; Tomasi</u>		X	
<u>FAST</u>		X	
<u>Laplacian of Gaussian</u>		X	X
<u>Difference of Gaussians</u>		X	X
<u>Determinant of Hessian</u>		X	X
<u>MSER</u>			X